

Motivation
oooooo

Background
ooooooo

Prox Algorithms
oooooooo

What if...
ooooooo

More adventuring
oooo

References
ooooo

A quest for theoretically-sound optimization


Zev Woodstock

Technische Universität Berlin & Zuse Institute Berlin

February 2024



“How did I get here?” -David Byrne



“How did I get here?” -David Byrne



Lise

6.05 – 8 petaflop/second
(roughly 75 – 100 IBM Watsons)



Theoretically-sound optimization

- 1.** Motivation
- 2.** Background: Theory vs practice
- 3.** Proximity operators: Algorithmic bells and whistles
- 4.** Splitting FW: What if the “usual” tools fail us?
- 5.** More adventuring

What is optimization?

Optimization in a nutshell ($\mathcal{H} = \mathbb{R}^n$ or any real Hilbert space)

- **Objective function** $f: \mathcal{H} \rightarrow \mathbb{R} \cup \{+\infty\}$.
e.g., data fidelity in ML, energy efficiency, profit, statistical error, ...
- An “optimal” $x \in \mathcal{H}$ makes $f(x)$ the smallest or largest
e.g., **minimize** error, **maximize** efficiency




image: towardsdatascience.com

$$\underset{x \in \mathcal{H}}{\text{minimize}} \quad f(x)$$


$$\iota_{\mathcal{C}}(x) = \begin{cases} 0 & \text{if } x \in \mathcal{C} \\ +\infty & \text{otherwise.} \end{cases}$$

Constraint set(s) $\mathcal{C} \subset \mathcal{H}$

e.g., \mathbb{R}_+^N , \mathbb{S}_+^N , hypercube, solution set of an inverse problem, ...


$$\underset{x \in \mathcal{C}}{\text{minimize}} \quad \tilde{f}(x) = \underset{x \in \mathcal{H}}{\text{minimize}} \quad \underbrace{\tilde{f}(x) + \iota_{\mathcal{C}}(x)}_f$$

Modeling via optimization



[Torelli et al., *Med. Phys.*, 2023]

image: [Fu et al., *Tech. Cancer Res. Treatment*, 2023]



[Sartori & Buriol, *Comput. Oper. Res.*, 2020]

Modeling via optimization



0	2	15	0	0	11	10	0	0	0	0	9	9	0	0	0	
0	0	0	4	60	152	236	255	255	255	95	61	32	0	0	29	
0	10	16	115	238	255	244	245	245	250	249	255	222	103	10	0	
0	14	170	255	255	244	254	255	253	245	255	249	253	251	124	1	
2	90	255	228	255	255	251	254	211	141	116	127	215	251	238	255	49
13	217	243	255	155	33	226	58	2	0	10	19	232	255	255	36	
16	229	252	254	49	12	0	7	7	0	70	237	252	255	62		
0	14	245	255	212	25	11	9	3	0	115	236	243	255	137	0	
0	8	252	250	248	215	60	0	3	20	252	255	248	144	6	0	
0	13	113	255	255	245	255	182	181	248	257	242	208	36	0	19	
1	0	5	111	251	255	241	255	247	255	241	162	17	0	7	0	
0	0	0	4	16	251	259	246	254	253	250	120	11	0	1	0	
0	0	4	91	255	255	255	248	252	255	244	255	182	10	0	4	
0	22	206	252	246	251	241	100	24	110	255	245	194	9	0		
0	111	255	242	255	158	24	0	0	6	30	255	232	230	56	0	
0	218	251	250	137	7	11	0	0	0	2	62	255	250	175	3	
0	173	255	255	103	9	20	0	13	3	12	182	251	245	61	0	
0	107	251	241	255	230	98	55	19	118	217	248	253	255	52	4	
0	18	146	250	255	247	255	255	249	255	240	255	129	0	5		
0	0	23	115	215	255	250	248	255	255	248	248	118	14	12	0	
0	0	6	1	0	52	151	233	255	257	141	37	0	0	4	1	
0	0	5	5	0	0	0	0	0	14	1	0	6	6	0	0	

0 2 15 0 0 11 10 0 0 0 0 9 9 0 0 0
0 0 0 4 60 152 236 255 255 177 95 61 32 0 0 29
0 10 16 115 238 255 244 245 245 250 249 255 222 103 10 0
0 14 170 255 255 244 254 254 253 245 255 249 253 251 124 1
2 98 255 228 255 255 251 254 211 141 116 127 215 251 238 255 49
13 217 243 255 155 33 226 58 2 0 10 19 232 255 255 36
16 229 252 254 49 12 0 0 7 7 0 70 237 252 255 62
6 141 245 255 212 25 11 9 3 0 115 236 243 255 137 0
0 87 252 250 248 215 60 0 3 20 252 255 248 144 6 0
0 13 113 255 255 245 255 182 181 248 257 242 208 36 0 19
1 0 5 117 251 255 241 255 247 247 255 241 161 17 0 7 0
0 0 0 4 58 251 255 246 254 253 256 120 11 0 1 0
0 0 4 97 255 255 255 248 252 255 244 255 182 10 0 4
0 22 206 252 246 251 241 100 24 113 255 245 255 194 9 0
0 111 255 242 255 158 24 0 0 6 39 255 232 230 56 0
0 218 251 250 137 7 11 0 0 0 2 62 255 250 125 3
0 173 255 255 101 9 20 0 13 3 13 182 251 245 61 0
0 107 251 241 255 230 98 55 19 118 217 248 253 255 52 4
0 18 146 250 255 247 255 255 249 255 240 255 129 0 5
0 0 23 113 215 255 250 248 255 255 248 248 118 14 12 0
0 0 6 1 0 52 153 233 255 257 141 37 0 0 4 1
0 0 5 5 0 0 0 0 0 14 1 0 6 6 0 0 0

image: towardsdatascience.com


Modeling via optimization



0	2	15	0	0	11	10	0	0	0	0	9	9	0	0	0
0	0	0	4	60	152	236	255	255	377	95	61	32	0	0	29
0	10	16	111	238	255	244	245	245	250	249	255	222	103	10	0
0	14	170	255	255	244	254	255	253	245	255	249	253	251	124	1
2	90	255	228	255	255	251	254	211	141	116	127	215	251	238	49
13	217	243	255	155	33	226	58	2	0	10	19	232	255	255	36
16	229	252	254	49	12	0	7	7	0	70	237	252	255	62	
6	141	245	255	212	25	11	9	3	0	115	236	243	255	137	0
0	87	252	250	248	215	60	0	3	20	252	255	248	144	6	0
0	13	113	255	255	245	255	182	181	248	252	242	208	36	0	19
1	0	5	117	251	255	241	255	247	255	241	165	17	0	7	0
0	0	0	4	16	251	255	246	254	253	250	220	11	0	1	0
0	0	4	91	255	255	255	248	252	255	244	255	182	10	0	4
0	22	206	252	246	251	241	100	24	110	255	245	194	9	0	
0	111	255	242	255	158	24	0	0	6	35	255	232	230	56	0
0	218	251	250	137	7	11	0	0	0	2	62	255	250	125	3
0	173	255	255	103	9	20	0	13	3	13	182	251	245	61	0
0	107	251	241	255	230	98	55	19	118	217	248	253	255	52	4
0	18	146	250	255	247	255	255	249	255	240	255	129	0	5	
0	0	23	113	215	255	250	248	255	255	248	248	118	14	12	0
0	0	6	1	0	52	151	233	255	257	141	37	0	0	4	1
0	0	5	5	0	0	0	0	0	14	1	0	6	6	0	0

0 2 15 0 0 11 10 0 0 0 0 9 9 0 0 0
0 0 0 4 60 152 236 255 255 377 95 61 32 0 0 29
0 10 16 111 238 255 244 245 245 250 249 255 222 103 10 0
0 14 170 255 255 244 254 255 253 245 255 249 253 251 124 1
2 90 255 228 255 255 251 254 211 141 116 127 215 251 238 49
13 217 243 255 155 33 226 58 2 0 10 19 232 255 255 36
16 229 252 254 49 12 0 0 7 7 0 70 237 252 255 62
6 141 245 255 212 25 11 9 3 0 115 236 243 255 137 0
0 87 252 250 248 215 60 0 3 20 252 255 248 144 6 0
0 13 113 255 255 245 255 182 181 248 252 242 208 36 0 19
1 0 5 117 251 255 241 255 247 255 241 165 17 0 7 0
0 0 0 4 16 251 255 246 254 253 250 220 11 0 1 0
0 0 4 91 255 255 255 248 252 255 244 255 182 10 0 4
0 22 206 252 246 251 241 100 24 110 255 245 194 9 0
0 111 255 242 255 158 24 0 0 6 35 255 232 230 56 0
0 218 251 250 137 7 11 0 0 0 2 62 255 250 125 3
0 173 255 255 103 9 20 0 13 3 13 182 251 245 61 0
0 107 251 241 255 230 98 55 19 118 217 248 253 255 52 4
0 18 146 250 255 247 255 255 249 255 240 255 129 0 5
0 0 23 113 215 255 250 248 255 255 248 248 118 14 12 0
0 0 6 1 0 52 151 233 255 257 141 37 0 0 4 1
0 0 5 5 0 0 0 0 0 14 1 0 6 6 0 0

image: towardsdatascience.com



Modeling via optimization

he actions are few and p
efined to a larger colle
ctional to retain the str
einement. Because of t
he level will greatly infl
re is insufficient infor
e should decide as littl
be made in an arbitrary

Original



Observation

he actions and few and p
efined to a larger colle
ctional to retain the str
einement. Because of t
he level will greatly infl
re is insufficient infor
e should decide as littl
be made in an arbitrary

Recovery

Modeling via optimization

he actions are few and p
efined to a larger colle
ctional to retain the str
einement. Because of t
he level will greatly infl
re is insufficient infor
e should decide as littl
be made in an arbitrary

Original



Observation

he actions and few and p
efined to a larger colle
ctional to retain the str
einement. Because of t
he level will greatly infl
re is insufficient infor
e should decide as littl
be made in an arbitrary

Recovery




Original



Observation



Recovered stars



Recovered galaxy

Some fundamental questions



- What are the roadblocks to *provably* solving optimization problems?
 - Nonconvexity, nonsmoothness, and bears – oh my!

Some fundamental questions



- What are the roadblocks to *provably* solving optimization problems?
 - Nonconvexity, nonsmoothness, and bears – oh my!
- What *theoretically-sound* algorithms exist?

Some fundamental questions



- What are the roadblocks to *provably* solving optimization problems?
 - Nonconvexity, nonsmoothness, and bears – oh my!
- What *theoretically-sound* algorithms exist, and can we do better?
 - Splitting, Parallelization, Extrapolation, Asynchronous computation

Motivation
oooooo

Background
•ooooo

Prox Algorithms
oooooooo

What if...
ooooooo

More adventuring
oooo

References
ooooo

Theoretically-sound optimization

- 1.** Motivation
- 2.** Background: Theory vs practice
- 3.** Proximity operators: Algorithmic bells and whistles
- 4.** Splitting FW: What if the “usual” tools fail us?
- 5.** More adventuring

Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

$\text{Argmin } f = \{x \in \mathcal{H} \mid f(x) = \inf f(\mathcal{H})\}$ is the set of minimizers of f .

Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

$\text{Argmin } f = \{x \in \mathcal{H} \mid f(x) = \inf f(\mathcal{H})\}$ is the **set of minimizers** of f .


f is **convex** if, for all $x, y \in \mathcal{H}$ and $\alpha \in (0, 1)$, $f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$

Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

$\text{Argmin } f = \{x \in \mathcal{H} \mid f(x) = \inf f(\mathcal{H})\}$ is the **set of minimizers** of f .

f is **convex** if, for all $x, y \in \mathcal{H}$ and $\alpha \in (0, 1)$, $f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$




Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

$\text{Argmin } f = \{x \in \mathcal{H} \mid f(x) = \inf f(\mathcal{H})\}$ is the **set of minimizers** of f .

f is **convex** if, for all $x, y \in \mathcal{H}$ and $\alpha \in (0, 1)$, $f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$





Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

$\text{Argmin } f = \{x \in \mathcal{H} \mid f(x) = \inf f(\mathcal{H})\}$ is the **set of minimizers** of f .

f is **convex** if, for all $x, y \in \mathcal{H}$ and $\alpha \in (0, 1)$, $f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$




Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

$\text{Argmin } f = \{x \in \mathcal{H} \mid f(x) = \inf f(\mathcal{H})\}$ is the **set of minimizers** of f .

C is **convex** if, for all $x, y \in C$ and $\alpha \in (0, 1)$, $\alpha x + (1 - \alpha)y \in C$.



Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

$\text{Argmin } f = \{x \in \mathcal{H} \mid f(x) = \inf f(\mathcal{H})\}$ is the **set of minimizers** of f .

f is **convex** if, for all $x, y \in \mathcal{H}$ and $\alpha \in (0, 1)$, $f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$

Let $\Gamma_0(\mathcal{H}) = \{f : \mathcal{H} \rightarrow \mathbb{R} \cup \{+\infty\} \mid f \text{ is convex, lower-semicontinuous, and proper}\}$,

Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

$\text{Argmin } f = \{x \in \mathcal{H} \mid f(x) = \inf f(\mathcal{H})\}$ is the **set of minimizers** of f .

f is **convex** if, for all $x, y \in \mathcal{H}$ and $\alpha \in (0, 1)$, $f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$

Let $\Gamma_0(\mathcal{H}) = \{f : \mathcal{H} \rightarrow \mathbb{R} \cup \{+\infty\} \mid f \text{ is convex, lower-semicontinuous, and proper}\}$,
e.g., e^x , $-\ln(x)$, $\|\cdot\|^2$, ReLU, Hinge loss, $\|Ax + b\|$, $\|\cdot\|_1$, ι_C (C convex and closed),

Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

$\text{Argmin } f = \{x \in \mathcal{H} \mid f(x) = \inf f(\mathcal{H})\}$ is the **set of minimizers** of f .

f is **convex** if, for all $x, y \in \mathcal{H}$ and $\alpha \in (0, 1)$, $f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$

Let $\Gamma_0(\mathcal{H}) = \{f : \mathcal{H} \rightarrow \mathbb{R} \cup \{+\infty\} \mid f \text{ is convex, lower-semicontinuous, and proper}\}$,

e.g., e^x , $-\ln(x)$, $\|\cdot\|^2$, ReLU, Hinge loss, $\|Ax + b\|$, $\|\cdot\|_1$, ι_C (C convex and closed),
 $\sup\{f_i \mid i \in I\}$, affine composition, positive linear combinations, ...

Our setting

\mathcal{H} is a real Hilbert space with inner product $\langle \cdot | \cdot \rangle$,
e.g., \mathbb{R}^n with the dot product $\langle x | y \rangle = x^T y$.

$\text{Argmin } f = \{x \in \mathcal{H} \mid f(x) = \inf f(\mathcal{H})\}$ is the **set of minimizers** of f .

f is **convex** if, for all $x, y \in \mathcal{H}$ and $\alpha \in (0, 1)$, $f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$


Let $\Gamma_0(\mathcal{H}) = \{f : \mathcal{H} \rightarrow \mathbb{R} \cup \{+\infty\} \mid f \text{ is convex, lower-semicontinuous, and proper}\}$,

e.g., e^x , $-\ln(x)$, $\|\cdot\|^2$, ReLU, Hinge loss, $\|Ax + b\|$, $\|\cdot\|_1$, ι_C (C convex and closed),
 $\sup\{f_i \mid i \in I\}$, affine composition, positive linear combinations, ...
not $\|\mathcal{N}(x) - d\|$ for multilayer neural networks

“Traditioooon”

-Tevye, Fiddler on the Roof

f is **L -smooth** ($L \geq 0$) if it is differentiable and $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$.



“Traditioooon” -Tevye, Fiddler on the Roof


f is **L -smooth** ($L \geq 0$) if it is differentiable and $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$.

Gradient Descent

Let $f: \mathcal{H} \rightarrow \mathbb{R}$ be L -smooth and suppose $\text{Argmin } f \neq \emptyset$. Let $x_0 \in \mathcal{H}$, $\varepsilon > 0$ and for every $n \in \mathbb{N}$, set

$$x_{n+1} = x_n - \lambda_n \nabla f(x_n), \quad \text{where } \lambda_n \in \left[\varepsilon, \frac{2}{L} - \varepsilon \right] \quad (\text{GD})$$

If $f \in \Gamma_0(\mathcal{H})$, then $(x_n)_{n \in \mathbb{N}}$ converges to a minimizer of f .



Foe #1: Non-convexity

f is **L -smooth** ($L \geq 0$) if it is differentiable and $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$.


Gradient Descent

Let $f: \mathcal{H} \rightarrow \mathbb{R}$ be L -smooth and suppose $\text{Argmin } f \neq \emptyset$. Let $x_0 \in \mathcal{H}$, $\varepsilon > 0$ and for every $n \in \mathbb{N}$, set

$$x_{n+1} = x_n - \lambda_n \nabla f(x_n), \quad \text{where } \lambda_n \in \left[\varepsilon, \frac{2}{L} - \varepsilon \right] \quad (\text{GD})$$

If $f \notin \Gamma_0(\mathcal{H})$, then $(x_n)_{n \in \mathbb{N}}$ converges to a $\underbrace{\text{stationary point}}_{(\nabla f(x^*) = 0)}$.

For x_0 sufficiently close to a minimizer, $(x_n)_{n \in \mathbb{N}}$ converges to one.



“Traditioooon” -Tevye, Fiddler on the Roof

f is **L -smooth** ($L \geq 0$) if it is differentiable and $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$.


$$f = \sum_{1 \leq i \leq m} f_i$$

Gradient Descent

Let $f: \mathcal{H} \rightarrow \mathbb{R}$ be L -smooth and suppose $\text{Argmin } f \neq \emptyset$. Let $x_0 \in \mathcal{H}$, $\varepsilon > 0$ and for every $n \in \mathbb{N}$, set

$$x_{n+1} = x_n - \lambda_n \nabla f(x_n), \quad \text{where } \lambda_n \in \left[\varepsilon, \frac{2}{L} - \varepsilon \right] \quad (\text{GD})$$

If $f \in \Gamma_0(\mathcal{H})$, then $(x_n)_{n \in \mathbb{N}}$ converges to a minimizer of f .



“Traditioooon” -Tevye, Fiddler on the Roof

f is **L -smooth** ($L \geq 0$) if it is differentiable and $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$.


$$f = \sum_{1 \leq i \leq m} f_i$$

Stochastic Gradient Descent (one variant)

Let $f: \mathcal{H} \rightarrow \mathbb{R}$ be L -smooth and suppose $\text{Argmin } f \neq \emptyset$. Let $x_0 \in \mathcal{H}$ and for every $n \in \mathbb{N}$, set

$$x_{n+1} = x_n - \frac{1}{n+1} \nabla f_{i_n}(x_n), \quad \text{where } i_n \sim U(\{1, \dots, m\}) \quad (\text{SGD})$$

If $f \in \Gamma_0(\mathcal{H})$, then $\mathbb{E}[f(x_n)]$ converges to $\inf_{x \in \mathcal{H}} f(x)$.



Why can't we take the eagles to Mordor?

(A reasonable question to ask, if we didn't read the books)

A common paradigm:

1. Define an **objective function**
2. Optimize with an efficient algorithm, e.g., SGD with **algorithmic differentiation (AD)**

Why can't we take the eagles to Mordor?

(A reasonable question to ask, if we didn't read the books)

A common paradigm:

1. Define an **objective function**
2. Optimize with an efficient algorithm, e.g., SGD with **algorithmic differentiation (AD)**

[Pontil et al., *Numer. Algorithms*, 2019]

Training a sparse linear binary classifier




image: adeveloperdiary.com

Why can't we take the eagles to Mordor?

(A reasonable question to ask, if we didn't read the books)

A common paradigm:

1. Define an **objective function**
2. Optimize with an efficient algorithm, e.g., SGD with **algorithmic differentiation (AD)**

[Pontil et al., *Numer. Algorithms*, 2019]

Training a sparse linear binary classifier




image: adeveloperdiary.com

$$\begin{aligned} \underset{x \in \mathbb{R}^n}{\text{minimize}} \quad & \sum_{i \in I_1} \max\{0, 1 - \langle x | a_i \rangle\} + \\ & \sum_{i \in I_2} \max\{0, 1 + \langle x | a_i \rangle\} + \lambda \|x\|_1 \end{aligned}$$

Foe #2: Non-differentiability

A common paradigm:

1. Define an **objective function**
2. Optimize with an efficient algorithm, e.g., SGD with **algorithmic differentiation** (AD)

Issue: For many **objective functions**, a gradient does not exist.

[Pontil et al., *Numer. Algorithms*, 2019]

Training a sparse linear binary classifier




image: adeveloperdiary.com

$$\begin{aligned} \text{minimize}_{x \in \mathbb{R}^n} & \sum_{i \in I_1} \max\{0, 1 - \langle x | a_i \rangle\} + \\ & \sum_{i \in I_2} \max\{0, 1 + \langle x | a_i \rangle\} + \lambda \|x\|_1 \end{aligned}$$

Foe #2: Non-differentiability

A common paradigm:

1. Define an **objective function**
2. Optimize with an efficient algorithm, e.g., SGD with **algorithmic differentiation** (AD)

Issue: For many **objective functions**, a **gradient** does not exist.

Engineers* :



image: ripley's.com


*:Some mathematicians at heart exceptions exist

How do we solve $\nabla f = 0$ when ∇f doesn't exist?

How do we solve $\nabla f = 0$ when ∇f doesn't exist?

If $f \in \Gamma_0(\mathcal{H})$ is differentiable at $x \in \mathcal{H}$, then


$$(\forall y \in \mathcal{H}) \quad \left\langle y - x \mid \nabla f(x) \right\rangle + f(x) \leq f(y).$$



How do we solve $\nabla f = 0$ when ∇f doesn't exist?

A **subgradient** $g \in \mathcal{H}$ of $f: \mathcal{H} \rightarrow]-\infty, +\infty]$ at $x \in \mathcal{H}$ satisfies

$$(\forall y \in \mathcal{H}) \quad \langle y - x \mid g \rangle + f(x) \leq f(y).$$




How do we solve $\nabla f = 0$ when ∇f doesn't exist?

A **subgradient** $g \in \mathcal{H}$ of $f: \mathcal{H} \rightarrow]-\infty, +\infty]$ at $x \in \mathcal{H}$ satisfies

$$(\forall y \in \mathcal{H}) \quad \langle y - x \mid g \rangle + f(x) \leq f(y).$$

The **subdifferential** $\partial f(x) \subset \mathcal{H}$ is the set containing **all subgradients** of f at x .





How do we solve $\nabla f = 0$ when ∇f doesn't exist?

A **subgradient** $g \in \mathcal{H}$ of $f: \mathcal{H} \rightarrow]-\infty, +\infty]$ at $x \in \mathcal{H}$ satisfies

$$(\forall y \in \mathcal{H}) \quad \langle y - x \mid g \rangle + f(x) \leq f(y).$$

The **subdifferential** $\partial f(x) \subset \mathcal{H}$ is the set containing **all subgradients** of f at x .

Example: $f = |\cdot|$: What do we do at zero?





How do we solve $\nabla f = 0$ when ∇f doesn't exist?

A **subgradient** $g \in \mathcal{H}$ of $f: \mathcal{H} \rightarrow]-\infty, +\infty]$ at $x \in \mathcal{H}$ satisfies

$$(\forall y \in \mathcal{H}) \quad \langle y - x \mid g \rangle + f(x) \leq f(y).$$

The **subdifferential** $\partial f(x) \subset \mathcal{H}$ is the set containing **all subgradients** of f at x .

Example: $f = |\cdot|$: What do we do at zero? $\partial f(0) = [-1, 1]$



How do we solve $\nabla f = 0$ when ∇f doesn't exist?

A **subgradient** $g \in \mathcal{H}$ of $f: \mathcal{H} \rightarrow]-\infty, +\infty]$ at $x \in \mathcal{H}$ satisfies

$$(\forall y \in \mathcal{H}) \quad \langle y - x \mid g \rangle + f(x) \leq f(y).$$

The **subdifferential** $\partial f(x) \subset \mathcal{H}$ is the set containing **all subgradients** of f at x .

Fermat's Rule

Let $x \in \mathcal{H}$. Then $x \in \text{Argmin } f$ if and only if $0 \in \partial f(x)$.

How do we solve $\nabla f = 0$ when ∇f doesn't exist?

A **subgradient** $g \in \mathcal{H}$ of $f: \mathcal{H} \rightarrow]-\infty, +\infty]$ at $x \in \mathcal{H}$ satisfies

$$(\forall y \in \mathcal{H}) \quad \langle y - x | g \rangle + f(x) \leq f(y).$$

The **subdifferential** $\partial f(x) \subset \mathcal{H}$ is the set containing **all subgradients** of f at x .

Fermat's Rule

Let $x \in \mathcal{H}$. Then $x \in \text{Argmin } f$ if and only if $0 \in \partial f(x)$.

Proof:

$$\begin{aligned} 0 \in \partial f(x) &\Leftrightarrow (\forall y \in \mathcal{H}) \quad \langle y - x | 0 \rangle + f(x) \leq f(y) \\ &\Leftrightarrow (\forall y \in \mathcal{H}) \quad f(x) \leq f(y) \\ &\Leftrightarrow x \in \text{Argmin } f \end{aligned}$$

How do we solve $\nabla f = 0$ when ∇f doesn't exist?

A **subgradient** $g \in \mathcal{H}$ of $f: \mathcal{H} \rightarrow]-\infty, +\infty]$ at $x \in \mathcal{H}$ satisfies

$$(\forall y \in \mathcal{H}) \quad \langle y - x \mid g \rangle + f(x) \leq f(y).$$

The **subdifferential** $\partial f(x) \subset \mathcal{H}$ is the set containing **all subgradients** of f at x .

∂f is useful for developing **both** $\underbrace{\text{optimality criterion and}}_{0 \in \partial f(x)}$ algorithms.

Goal: “ $0 \in \partial f(x)$ ”. Which path do we take?



image: centralldm.es

Goal: “ $0 \in \partial f(x)$ ”. Which path do we take?

Some *provenly-convergent* (first-order) algorithm classes:

- Subgradient-projections (e.g., in [C. & ZW, *IEEE EUSIPCO*, 2020])
- Proximity operators (e.g., in [C., B., & ZW, *IEEE ICASSP*, 2022], [C. & ZW, *J. Approx. Theory*, 2021], [C. & ZW, *SIAM J. Imaging Sci.*, 2023])
- Conditional Gradient / “Frank-Wolfe” (e.g., in [ZW & P., 2024], [K., P., W., & ZW, *Opt. Methods. Softw.*, 2024])
- Abs-smooth Optimization (e.g., [K., P., W., & ZW, *Opt. Methods. Softw.*, 2024])
- Bundle methods, Barrier methods, Lagrangian methods, ...

Goal: “ $0 \in \partial f(x)$ ”. Which path do we take?

Some *provenly-convergent* (first-order) algorithm classes:

- Subgradient-projections (e.g., in [C. & ZW, *IEEE EUSIPCO*, 2020])
- **Proximity operators** (e.g., in [C., B., & ZW, *IEEE ICASSP*, 2022], [C. & ZW, *J. Approx. Theory*, 2021], [C. & ZW, *SIAM J. Imaging Sci.*, 2023])
- **Conditional Gradient / “Frank-Wolfe”** (e.g., in [ZW & P., 2024], [K., P., W., & ZW, *Opt. Methods. Softw.*, 2024])
- Abs-smooth Optimization (e.g., [K., P., W., & ZW, *Opt. Methods. Softw.*, 2024])
- Bundle methods, Barrier methods, Lagrangian methods, ...


Theoretically-sound optimization

- 1.** Motivation
- 2.** Background: Theory vs practice
- 3.** Proximity operators: Algorithmic bells and whistles
- 4.** Splitting FW: What if the “usual” tools fail us?
- 5.** More adventuring

Proximity operators: a new hope

The **proximity operator** of f at $x \in \mathcal{H}$ is

$$\text{prox}_f(x) = \underset{u \in \mathcal{H}}{\operatorname{Argmin}} f(u) + \frac{1}{2} \|x - u\|^2$$




Proximity operators: a new hope

The **proximity operator** of f at $x \in \mathcal{H}$ is

$$\text{prox}_f(x) = \underset{u \in \mathcal{H}}{\operatorname{Argmin}} f(u) + \frac{1}{2} \|x - u\|^2$$


→ For $f \in \Gamma_0(\mathcal{H})$ and $x \in \mathcal{H}$, $\text{prox}_f(x)$ is unique.



Proximity operators: a new hope

The **proximity operator** of f at $x \in \mathcal{H}$ is

$$\text{prox}_f(x) = \underset{u \in \mathcal{H}}{\operatorname{Argmin}} f(u) + \frac{1}{2} \|x - u\|^2$$




→ For $f \in \Gamma_0(\mathcal{H})$ and $x \in \mathcal{H}$, $\text{prox}_f(x)$ is unique. Defines an operator $\text{prox}_f: \mathcal{H} \rightarrow \mathcal{H}$.

Proximity operators: a new hope

The **proximity operator** of f at $x \in \mathcal{H}$ is

$$\text{prox}_f(x) = \underset{u \in \mathcal{H}}{\operatorname{Argmin}} f(u) + \frac{1}{2} \|x - u\|^2$$




- For $f \in \Gamma_0(\mathcal{H})$ and $x \in \mathcal{H}$, $\text{prox}_f(x)$ is unique. Defines an operator $\text{prox}_f: \mathcal{H} \rightarrow \mathcal{H}$.
- Projection onto **closed convex set C** : $\text{prox}_{\iota_C}(x) = \underset{u \in C}{\operatorname{Argmin}} \|x - u\|^2 = \text{proj}_C x$.

Proximity operators: a new hope

The **proximity operator** of f at $x \in \mathcal{H}$ is

$$\text{prox}_f(x) = \underset{u \in \mathcal{H}}{\operatorname{Argmin}} f(u) + \frac{1}{2} \|x - u\|^2$$



- For $f \in \Gamma_0(\mathcal{H})$ and $x \in \mathcal{H}$, $\text{prox}_f(x)$ is unique. Defines an operator $\text{prox}_f: \mathcal{H} \rightarrow \mathcal{H}$.
- Projection onto **closed convex set C** : $\text{prox}_{\iota_C}(x) = \underset{u \in C}{\operatorname{Argmin}} \|x - u\|^2 = \text{proj}_C x$.
- [Martinet, *Fr. Inf. Rech. Oper.*, 1970] (translated / modernized):
Let $f \in \Gamma_0(\mathcal{H})$ be such that $\operatorname{Argmin} f \neq \emptyset$. Let $\gamma > 0$, $x_0 \in \mathcal{H}$, and set


$$x_{n+1} = \text{prox}_{\gamma f} x_n.$$

Then $(x_n)_{n \in \mathbb{N}} \rightharpoonup x^* \in \operatorname{Argmin} f$.

Proximity operators: a new hope

The **proximity operator** of f at $x \in \mathcal{H}$ is

$$\text{prox}_f(x) = \underset{u \in \mathcal{H}}{\operatorname{Argmin}} f(u) + \frac{1}{2} \|x - u\|^2$$



- For $f \in \Gamma_0(\mathcal{H})$ and $x \in \mathcal{H}$, $\text{prox}_f(x)$ is unique. Defines an operator $\text{prox}_f: \mathcal{H} \rightarrow \mathcal{H}$.
- Projection onto **closed convex set C** : $\text{prox}_{\iota_C}(x) = \underset{u \in C}{\operatorname{Argmin}} \|x - u\|^2 = \text{proj}_C x$.
- [Martinet, *Fr. Inf. Rech. Oper.*, 1970] (translated / modernized):
Let $f \in \Gamma_0(\mathcal{H})$ be such that $\operatorname{Argmin} f \neq \emptyset$. Let $\gamma > 0$, $x_0 \in \mathcal{H}$, and set

$$x_{n+1} = \text{prox}_{\gamma f} x_n.$$

Then $(x_n)_{n \in \mathbb{N}} \rightharpoonup x^* \in \operatorname{Argmin} f$. **Issue:** prox_f might be hard to compute.

Evolution of prox-based algorithms

$$f = \sum_{1 \leq i \leq m} f_i$$

$(\text{prox}_{f_i})_{1 \leq i \leq m}$ are simpler than prox_f .

Evolution of prox-based algorithms

$$f = \sum_{1 \leq i \leq m} f_i$$

$(\text{prox}_{f_i})_{1 \leq i \leq m}$ are simpler than prox_f .

Open-source repo's:

Python/Matlab:

[proximity-operator.net](#),

Julia:

[ProximalOperators.jl](#) (Github)

Evolution of prox-based algorithms

$$f = \sum_{1 \leq i \leq m} f_i$$

$(\text{prox}_{f_i})_{1 \leq i \leq m}$ are simpler than prox_f .

Algorithmic “Splitting” mentality:



prox_f



$\text{prox}_{f_1}, \dots, \text{prox}_{f_m}$

Evolution of prox-based algorithms

$$f = \sum_{1 \leq i \leq m} f_i$$

$(\text{prox}_{f_i})_{1 \leq i \leq m}$ are simpler than prox_f .

Algorithmic “Splitting” mentality:



prox_f



$\text{prox}_{f_1}, \dots, \text{prox}_{f_m}$

Pseudocode for many proximal splitting algorithms*

Require: Point $x_0 \in \mathcal{H}$, objective function f

```
1: for  $n = 0, 1$  to ... do
2:   # Preprocess  $x_n$ 
3:   for  $i = 1$  to  $m$  do
4:      $y_{i,n+1} \leftarrow$  evaluation of  $\text{prox}_{f_i}(\cdot)$ 
5:   end for
6:    $x_{n+1} \leftarrow$  combine  $(y_{i,n+1})_{1 \leq i \leq m}$ 
7: end for
```

*:e.g., Douglas-Rachford, ADMM, Chambolle-Pock, Forward-backward, Augmented Lagrangian, ...

Evolution of prox-based algorithms

$$f = \sum_{1 \leq i \leq m} f_i$$




image: baeldung.com

Pseudocode for many proximal splitting algorithms*

Require: Point $x_0 \in \mathcal{H}$, objective function f

```
1: for  $n = 0, 1$  to ... do
2:   # Preprocess  $x_n$ 
3:   for  $i = 1$  to  $m$  do
4:      $y_{i,n+1} \leftarrow$  evaluation of  $\text{prox}_{f_i}(\cdot)$ 
5:   end for
6:    $x_{n+1} \leftarrow$  combine  $(y_{i,n+1})_{1 \leq i \leq m}$ 
7: end for
```

*:e.g., Douglas-Rachford, ADMM, Chambolle-Pock, Forward-backward, Augmented Lagrangian, ...

Evolution of prox-based algorithms

$$f = \sum_{1 \leq i \leq m} f_i$$




image: baeldung.com

Pseudocode for many proximal splitting algorithms*

Require: Point $x_0 \in \mathcal{H}$, objective function f

```
1: for  $n = 0, 1$  to ... do
2:   # Preprocess  $x_n$ 
3:   for  $i = 1$  to  $m$  do
4:      $y_{i,n+1} \leftarrow$  evaluation of  $\text{prox}_{f_i}(\cdot)$ 
5:   end for
6:    $x_{n+1} \leftarrow$  combine  $(y_{i,n+1})_{1 \leq i \leq m}$ 
7: end for
```

*:e.g., Douglas-Rachford, ADMM, Chambolle-Pock, Forward-backward, Augmented Lagrangian, ...

Evolution of prox-based algorithms

$$f = \sum_{1 \leq i \leq m} f_i$$





image: baeldung.com

Pseudocode for block-iterative proximal splitting algorithms

Require: Point $x_0 \in \mathcal{H}$, objective function f

```
1: for  $n = 0, 1$  to ... do
2:   # Preprocess  $x_n$ ; Select  $I_n \subset \{1, \dots, m\}$ 
3:   for  $i = 1$  to  $m$  do
4:     if  $i \in I_n$  then
5:        $y_{i,n+1} \leftarrow$  evaluation of  $\text{prox}_{f_i}(\cdot)$ 
6:     else
7:        $y_{i,n+1} \leftarrow y_{i,n}$ 
8:     end if
9:   end for
10:   $x_{n+1} \leftarrow$  combine  $(y_{i,n+1})_{1 \leq i \leq m}$ 
11: end for
```

Block activation for image recovery ($m = 2$)



[C. & ZW, SIAM J. Imaging Sci., 2022] Relative error
for full-activation ($I_n = I$) versus block activation:

$$I_n = \begin{cases} \{1, 2\}, & \text{if } n \equiv 0 \pmod{5}; \\ \{2\}, & \text{if } n \not\equiv 0 \pmod{5}. \end{cases}$$

Evolution of prox-based algorithms

$$f = \sum_{1 \leq i \leq m} f_i$$




image: baeldung.com

Pseudocode for block-iterative proximal splitting algorithms*

Require: Point $x_0 \in \mathcal{H}$, objective function f

```
1: for  $n = 0, 1$  to ... do
2:   # Preprocess  $x_n$ ; Select  $I_n \subset \{1, \dots, m\}$ 
3:   for  $i = 1$  to  $m$  do
4:     if  $i \in I_n$  then
5:        $y_{i,n+1} \leftarrow$  evaluation of  $\text{prox}_{f_i}(\cdot)$ 
6:     else
7:        $y_{i,n+1} \leftarrow y_{i,n}$ 
8:     end if
9:   end for
10:   $x_{n+1} \leftarrow$  combine  $(y_{i,n+1})_{1 \leq i \leq m}$ 
11: end for
```

Evolution of prox-based algorithms

$$f = \sum_{1 \leq i \leq m} f_i$$

Two (currently separate) approaches:

- Asynchronous updates
- Extrapolated updates


Pseudocode for block-iterative proximal splitting algorithms*

Require: Point $x_0 \in \mathcal{H}$, objective function f

```
1: for  $n = 0, 1$  to ... do
2:   # Preprocess  $x_n$ ; Select  $I_n \subset \{1, \dots, m\}$ 
3:   for  $i = 1$  to  $m$  do
4:     if  $i \in I_n$  then
5:        $y_{i,n+1} \leftarrow$  evaluation of prox $f_i$ ( $\cdot$ )
6:     else
7:        $y_{i,n+1} \leftarrow y_{i,n}$ 
8:     end if
9:   end for
10:   $x_{n+1} \leftarrow$  combine  $(y_{i,n+1})_{1 \leq i \leq m}$ 
11: end for
```


Evolution of prox-based algorithms

Parallel and Synchronous




Evolution of prox-based algorithms

Parallel and Synchronous



Parallel and Asynchronous



Asynchrony: Projective Splitting Algorithms

[Eckstein & Svaiter, *Math Prog. A*, 2008]:

Coined “Projective splitting” (synchronous, not block-iterative).

[Combettes & Eckstein, *Math Prog. B*, 2018]:

Convergence proof with **asynchronous block-iterative** updates!

Asynchrony: Projective Splitting Algorithms

[Eckstein & Svaiter, *Math Prog. A*, 2008]:

Coined “Projective splitting” (synchronous, not block-iterative).

[Combettes & Eckstein, *Math Prog. B*, 2018]:

Convergence proof with **asynchronous block-iterative** updates!

[Combettes, Búi, & ZW, *IEEE ICASSP*, 2022]:

Numerical analysis (space and time complexity). For ML (training classifiers) and image processing, works better than other algorithms in its class.

Asynchrony: Projective Splitting Algorithms

[Eckstein & Svaiter, *Math Prog. A*, 2008]:

Coined “Projective splitting” (synchronous, not block-iterative).

[Combettes & Eckstein, *Math Prog. B*, 2018]:

Convergence proof with **asynchronous block-iterative** updates!

[Combettes, Búi, & ZW, *IEEE ICASSP*, 2022]:

Numerical analysis (space and time complexity). For ML (training classifiers) and image processing, works better than other algorithms in its class.

Only studied **synchronous** case!

Asynchrony: Projective Splitting Algorithms

[Eckstein & Svaiter, *Math Prog. A*, 2008]:

Coined “Projective splitting” (synchronous, not block-iterative).

[Combettes & Eckstein, *Math Prog. B*, 2018]:

Convergence proof with **asynchronous block-iterative** updates!

[Combettes, Búi, & ZW, *IEEE ICASSP*, 2022]:

Numerical analysis (space and time complexity). For ML (training classifiers) and image processing, works better than other algorithms in its class.

Only studied **synchronous** case!


Dua, Goel, Sharma, & ZW (ongoing work):

Analysis and experimentation for **asynchronous** case. AsyncProx.jl in development.

Over-relaxation (extrapolation)

Given a “descent direction” d (e.g., combination of $(y_{i,n+1})_{1 \leq i \leq m}$) from $x \in \mathcal{H}$,

$$x_+ = (1 - \alpha)x + \alpha(x + d)$$




Over-relaxation (extrapolation)

Given a “descent direction” d (e.g., combination of $(y_{i,n+1})_{1 \leq i \leq m}$) from $x \in \mathcal{H}$,

$$x_+ = (1 - \alpha)x + \alpha(x + d)$$

$0 \leq \alpha < 1$: under-relaxation

$1 < \alpha$: over-relaxation




Over-relaxation (extrapolation)

Given a “descent direction” d (e.g., combination of $(y_{i,n+1})_{1 \leq i \leq m}$) from $x \in \mathcal{H}$,

$$x_+ = (1 - \alpha)x + \alpha(x + d)$$

$0 \leq \alpha < 1$: under-relaxation

$1 < \alpha$: over-relaxation





$$\begin{aligned} x_{n+1} &= x_n - \alpha_n \frac{1}{L} \nabla f(x_n), & \text{where } \alpha_n \in [\varepsilon, 2 - \varepsilon] & \quad (\text{GD}) \\ &= (1 - \alpha_n)x_n + \alpha_n \left(x_n - \frac{1}{L} \nabla f(x_n) \right) \end{aligned}$$

Over-relaxation for fixed-point problems

[Combettes & ZW, *J. Approx. Theory*, 2021]:

A strongly-convergent algorithm with **affine-convex extrapolation**.

Example: EEG data (minimal-norm solution to an ill-posed nonlinear inverse problem;
recovery after 1000 iterations, < 1 minute.)



Theoretically-sound optimization

- 1.** Motivation
- 2.** Background: Theory vs practice
- 3.** Proximity operators: Algorithmic bells and whistles
- 4.** Splitting FW: What if the “usual” tools fail us?
- 5.** More adventuring

They stole my horse!

Splitting problem setup

Given a smooth function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ and compact convex sets $(C_i)_{1 \leq i \leq m}$

$$\text{minimize } f(x) \text{ subject to } x \in \bigcap_{1 \leq i \leq m} C_i, \quad (*)$$

Applications: data science, matrix decomposition, quantum computing, combinatorial graph theory

They stole my horse!

Splitting problem setup

Given a smooth function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ and compact convex sets $(C_i)_{1 \leq i \leq m}$

$$\text{minimize } f(x) \text{ subject to } x \in \bigcap_{1 \leq i \leq m} C_i, \quad (*)$$

Applications: data science, matrix decomposition, quantum computing, combinatorial graph theory

What if $(\text{proj}_{C_i})_{1 \leq i \leq m}$ are too expensive? \leftarrow e.g., in high-dimensional settings!

A spark of inspiration

Frank-Wolfe / “Conditional gradient” alg. [Naval Res. Logist. Quart., 1956]

Given a smooth function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ and a nonempty **compact convex set C** ,

$$\text{minimize } f(x) \text{ subject to } x \in C$$

Instead of projecting, use a *linear minimization oracle of C* ,

$$\text{LMO}_C: y \mapsto p \in \operatorname{Argmin}_{x \in C} \langle y | x \rangle \quad (\text{LMO})$$

$$x_{n+1} = x_n + \frac{1}{n+1} \left(\text{LMO}_C(\nabla f(x_n)) - x_n \right)$$



Marguerite Frank



Philip Wolfe

efficiency(LMOs) == efficiency(lembus bread)

[Combettes/Pokutta, '21]: For many constraints, C , proj_C is **more expensive** than LMO_C .
(e.g., nuclear norm ball, ℓ_1 ball, probability simplex, Birkhoff polytope, general LP, ...)

efficiency(LMOs) == efficiency(lembus bread)

[Combettes/Pokutta, '21]: For many constraints, C , proj_C is **more expensive** than LMO_C .
(e.g., nuclear norm ball, ℓ_1 ball, probability simplex, Birkhoff polytope, general LP, ...)

Example: Nuclear norm ball

For $x \in \mathbb{R}^{n \times n}$

$$\|x\|_{nuc} = \sum_{1 \leq i \leq n} \sigma_i(x).$$

For $\beta \geq 0$, $C = \{x \in \mathbb{R}^{n \times n} \mid \|x\|_{nuc} \leq \beta\}$,

efficiency(LMOs) == efficiency(lembus bread)

[Combettes/Pokutta, '21]: For many constraints, C , proj_C is **more expensive** than LMO_C .
(e.g., nuclear norm ball, ℓ_1 ball, probability simplex, Birkhoff polytope, general LP, ...)

Example: Nuclear norm ball

For $x \in \mathbb{R}^{n \times n}$

$$\|x\|_{nuc} = \sum_{1 \leq i \leq n} \sigma_i(x).$$

For $\beta \geq 0$, $C = \{x \in \mathbb{R}^{n \times n} \mid \|x\|_{nuc} \leq \beta\}$,

Full SVD

$n = 500$: 0.11 sec

$n = 1000$: 0.47 sec

$n = 2000$: 4.87 sec

$\text{proj}_C(x)$: requires a **full SVD!**

$(\sigma_1, \dots, \sigma_n, U, V)$, where $x = U \text{diag}(\sigma_1, \dots, \sigma_n) V^\top$

efficiency(LMOs) == efficiency(lembus bread)

[Combettes/Pokutta, '21]: For many constraints, C , proj_C is **more expensive** than LMO_C .
(e.g., nuclear norm ball, ℓ_1 ball, probability simplex, Birkhoff polytope, general LP, ...)

Example: Nuclear norm ball

For $x \in \mathbb{R}^{n \times n}$

$$\|x\|_{nuc} = \sum_{1 \leq i \leq n} \sigma_i(x).$$

For $\beta \geq 0$, $C = \{x \in \mathbb{R}^{n \times n} \mid \|x\|_{nuc} \leq \beta\}$,

$\text{proj}_C(x)$: requires a **full SVD!**

$(\sigma_1, \dots, \sigma_n, U, V)$, where $x = U \text{diag}(\sigma_1, \dots, \sigma_n) V^\top$

$\text{LMO}_C(x)$: requires only **first singular value/vectors**
 $(\sigma_1, U_1, V_1^\top)$

Full SVD

$n = 500$: 0.11 sec

$n = 1000$: 0.47 sec

$n = 2000$: 4.87 sec

Just $(\sigma_1, U_1, V_1^\top)$

$n = 500$: 0.0081 sec

$n = 1000$: 0.056 sec

$n = 2000$: 0.638 sec

Sure would be nice if we could use splitting...

$$\text{minimize } f(x) \text{ subject to } x \in \bigcap_{1 \leq i \leq m} C_i, \quad (*)$$

Sure would be nice if we could use splitting...

$$\text{minimize } f(x) \text{ subject to } x \in \bigcap_{1 \leq i \leq m} C_i, \quad (*)$$

Issue: Computing the LMO for $\bigcap_{1 \leq i \leq m} C_i$ is prohibitively expensive.

Sure would be nice if we could use splitting...

$$\text{minimize } f(x) \text{ subject to } x \in \bigcap_{1 \leq i \leq m} C_i, \quad (*)$$

Issue: Computing the LMO for $\bigcap_{1 \leq i \leq m} C_i$ is prohibitively expensive.

$(\text{LMO}_{C_i})_{1 \leq i \leq m}$ are easier to evaluate (e.g., repository: FrankWolfe.jl)

Sure would be nice if we could use splitting...

$$\text{minimize } f(x) \text{ subject to } x \in \bigcap_{1 \leq i \leq m} C_i, \quad (*)$$

Issue: Computing the LMO for $\bigcap_{1 \leq i \leq m} C_i$ is prohibitively expensive.

$(\text{LMO}_{C_i})_{1 \leq i \leq m}$ are easier to evaluate (e.g., repository: FrankWolfe.jl)

LMO-based *splitting algorithms*, enforce constraints via LMOs for the individual sets

Use $\text{LMO}_{C_1}, \text{LMO}_{C_2}, \dots$ instead of $\text{LMO}_{(\bigcap_{1 \leq i \leq m} C_i)}$

Sure would be nice if we could use splitting...

$$\text{minimize } f(x) \text{ subject to } x \in \bigcap_{1 \leq i \leq m} C_i, \quad (*)$$

Issue: Computing the LMO for $\bigcap_{1 \leq i \leq m} C_i$ is prohibitively expensive.

$(\text{LMO}_{C_i})_{1 \leq i \leq m}$ are easier to evaluate (e.g., repository: FrankWolfe.jl)

LMO-based *splitting algorithms*, enforce constraints via LMOs for the individual sets

Use $\text{LMO}_{C_1}, \text{LMO}_{C_2}, \dots$ instead of $\text{LMO}_{(\bigcap_{1 \leq i \leq m} C_i)}$

Relatively little has been done in this field.

Sure would be nice if we could use splitting...

$$\text{minimize } f(x) \text{ subject to } x \in \bigcap_{1 \leq i \leq m} C_i, \quad (*)$$

Issue: Computing the LMO for $\bigcap_{1 \leq i \leq m} C_i$ is prohibitively expensive.

$(\text{LMO}_{C_i})_{1 \leq i \leq m}$ are easier to evaluate (e.g., repository: FrankWolfe.jl)

LMO-based *splitting algorithms*, enforce constraints via LMOs for the individual sets

Use $\text{LMO}_{C_1}, \text{LMO}_{C_2}, \dots$ instead of $\text{LMO}_{(\bigcap_{1 \leq i \leq m} C_i)}$

Relatively little has been done in this field.

→ Unlike projections, LMOs are discontinuous.

Sure would be nice if we could use splitting...

$$\text{minimize } f(x) \text{ subject to } x \in \bigcap_{1 \leq i \leq m} C_i, \quad (*)$$

Issue: Computing the LMO for $\bigcap_{1 \leq i \leq m} C_i$ is prohibitively expensive.

$(\text{LMO}_{C_i})_{1 \leq i \leq m}$ are easier to evaluate (e.g., repository: FrankWolfe.jl)

LMO-based *splitting algorithms*, enforce constraints via LMOs for the individual sets

Use $\text{LMO}_{C_1}, \text{LMO}_{C_2}, \dots$ instead of $\text{LMO}_{(\bigcap_{1 \leq i \leq m} C_i)}$

Relatively little has been done in this field.

- Unlike projections, LMOs are discontinuous.
- “State-of-the-art” relies on inexact prox-based algorithms.

It's aliiive!

[ZW & P, 2024]: Split Conditional Gradient Algorithm

Require: Point $x_0 \in \frac{1}{m} \sum_{1 \leq i \leq m} C_i$, smooth function f

```
1: for  $t = 0, 1$  to ... do
2:   Choose penalty parameter  $\lambda_t \in ]0, +\infty[$ 
3:   Choose step size  $\gamma_t \in ]0, 1]$ 
4:    $g_t \leftarrow \nabla f(x_t)$ 
5:   for  $i = 1$  to  $m$  do
6:      $v_t^i \leftarrow \text{LMO}_{C_i}(g_t + \lambda_t(x_t^i - x_t))$ 
7:      $x_{t+1}^i \leftarrow x_t^i + \gamma_t(v_t^i - x_t^i)$ 
8:   end for
9:    $x_{t+1} \leftarrow \frac{1}{m} \sum_{1 \leq i \leq m} x_{t+1}^i$ 
10: end for
```

Practical advantages:

- Uses individual LMOs
- Lowest-known # LMO calls:
one LMO per set (per iteration)

It's aliiive!

[ZW & P, 2024]: Split Conditional Gradient Algorithm

Require: Point $x_0 \in \frac{1}{m} \sum_{1 \leq i \leq m} C_i$, smooth function f

```
1: for  $t = 0, 1$  to ... do
2:   Choose penalty parameter  $\lambda_t \in ]0, +\infty[$ 
3:   Choose step size  $\gamma_t \in ]0, 1]$ 
4:    $g_t \leftarrow \nabla f(x_t)$ 
5:   for  $i = 1$  to  $m$  do
6:      $v_t^i \leftarrow \text{LMO}_{C_i}(g_t + \lambda_t(x_t^i - x_t))$ 
7:      $x_{t+1}^i \leftarrow x_t^i + \gamma_t(v_t^i - x_t^i)$ 
8:   end for
9:    $x_{t+1} \leftarrow \frac{1}{m} \sum_{1 \leq i \leq m} x_{t+1}^i$ 
10:  end for
```

Practical advantages:

- Uses individual LMOs
- Lowest-known # LMO calls:
one LMO per set (per iteration)

Q: Does it actually solve (\star) ?

A: Yes.

$\gamma_t = \mathcal{O}(1/\sqrt{t})$ and
 $\lambda_t = \mathcal{O}(\ln t)$ work.
(whether or not f is convex).

Convergence

Theorem ([ZW & P., 2024])

Let f be L -smooth and let $(C_i)_{1 \leq i \leq m}$ be nonempty compact convex subsets of \mathcal{H} such that $\bigcap_{1 \leq i \leq m} C_i \neq \emptyset$. Let $\lambda_0 > 0$ and $\lambda_{t+1} = \lambda_t + (\sqrt{t} + 2)^{-2}$ and $\gamma_t = 2/(\sqrt{t} + 2)$. If f is **convex**, then

$$f\left(\frac{1}{m} \sum_{1 \leq i \leq m} \mathbf{x}_t^i\right) \rightarrow \inf_{x \in \bigcap_{1 \leq i \leq m} C_i} f(x)$$

and every accumulation point \mathbf{x}_∞ of $(\mathbf{x}_t)_{t \in \mathbb{N}}$ produces a solution

$$\frac{1}{m} \sum_{1 \leq i \leq m} \mathbf{x}_\infty^i \in \bigcap_{1 \leq i \leq m} C_i \text{ such that } f\left(\frac{1}{m} \sum_{1 \leq i \leq m} \mathbf{x}_\infty^i\right) = \inf_{x \in \bigcap_{1 \leq i \leq m} C_i} f(x).$$

Nonconvex convergence results too: arXiv:2311.05381


Theoretically-sound optimization

- 1.** Motivation
- 2.** Background: Theory vs practice
- 3.** Proximity operators: Algorithmic bells and whistles
- 4.** Splitting FW: What if the “usual” tools fail us?
- 5.** More adventuring


Abs-smooth optimization

Abs-smooth functions f include compositions of smooth functions, max, min, and $|\cdot|$

- Loss functions for multilayer **Neural Networks** with smooth and/or ReLU activation.
- Allows one to find a local minimizer on **non-convex** objective functions!
- Future work: Improve scalability.



Nonlinear inverse problems




Given $(p_i)_{1 \leq i \leq m}$, find $x^* \in \mathcal{H}$ such that

$$(\forall i \in \{1, \dots, m\}) \quad F_i(L_i x^*) = p_i$$

L_i are bounded linear operators, and
 $F_i \approx$ proximity operators.

→ To-do: Stability analysis



Outlook and future work

- Improved convergence rates and acceleration
- **Block-iterative** Frank-Wolfe algorithms
- Efficient prox/ LMO algorithms
- ...



Outlook and future work

- Improved convergence rates and acceleration
- **Block-iterative** Frank-Wolfe algorithms
- Efficient prox/ LMO algorithms
- ...

Potential collaborators: Hala Nelson, Minah Oh, Roger Thelwell, and more!

For students:

Proofs (MATH 245), sequences and series (236), gradients (237), linear algebra (238/300/434), optimization theory (340), coding experience (248/250/448/449), analysis (410/411).

REUs & Grad School in Berlin: iol.zib.de (Opt/ML), math-berlin.de



Motivation
oooooo

Background
oooooo

Prox Algorithms
oooooooo

What if...
ooooooo

More adventuring
oooo

References
●oooo

Thank you for your attention!

References

-  M. N. Búi, P. L. Combettes and ZW, **Block-activated algorithms for multicomponent fully nonsmooth minimization**
Proc. IEEE Int. Conf. Acoust. Speech Signal Process., pp 5428–5432, 2022.
-  P. L. Combettes, A. M. McDonald, C. A. Micchelli, and M. Pontil, **Learning with optimal interpolation norms**
Numer. Algorithms, vol. 81, no. 2, pp. 695–717, 2019.
-  P. L. Combettes and ZW, **Signal recovery from inconsistent nonlinear observations**
Proc. IEEE Int. Conf. Acoust. Speech Signal Process., pp 5872—5876, 2022.
-  P. L. Combettes and ZW, **A fixed point framework for recovering signals from nonlinear transformations**
Proc. IEEE Eur. Signal Proc., pp 2120–2124, 2020.
-  P. L. Combettes and ZW, **Reconstruction of functions from prescribed proximal points**
J. Approx. Theory, vol. 268, no. 105606, 2021

References

-  P. L. Combettes and ZW, A variational inequality model for the construction of signals from inconsistent nonlinear equations
SIAM J. Imaging Sci., vol. 15, no. 1, pp. 84–109, 2022
-  C. Combettes and S. Pokutta, Complexity of linear minimization and projection on some sets
Oper. Res. Lett., vol. 49, no. 4, pp. 565–571, 2021
-  M. Frank and P. Wolfe, An algorithm for quadratic programming
Naval Res. Logist. Quart., vol. 3, iss. 1–2, pp. 95–110, 1956
-  T. Kreimeier, S. Pokutta, A. Walther, and ZW, On a Frank-Wolfe approach for abs-smooth functions
Opt. Methods and Softw., DOI: 10.1080/10556788.2023.2296985, 2024
-  B. Martinet, régularisation d'inéquations variationnelles par approximations successives
Fr. Inf. Rech. Oper., Série rouge, 4 (R3):154–158, 1970.

References

-  C. S. Sartori and L. S. Buriol, A study on the pickup and delivery problem with time windows: Matheuristics and new instances
Comput. Oper. Res., vol. 124, art. 105065, 2020.
-  N. Torelli, D. Papp, and J. Unkelbach, Spatiotemporal fractionation schemes for stereotactic radiosurgery of multiple brain metastases
Med. Phys., vol. 50, no. 8, pp. 5095-5114, 2023.
-  ZW and S. Pokutta, [Splitting the conditional gradient algorithm](#)
arXiv:2311.05381, 2024

Supplement

$$f = \sum_{1 \leq i \leq m} f_i$$

$(\text{prox}_{f_i})_{1 \leq i \leq m}$ are simpler than prox_f .

e.g., Find $x \in \mathbb{S}_+^n \cap [\alpha, \beta]^{N \times N}$

$$\underset{x \in \mathbb{R}^n}{\text{minimize}} \quad \iota_{\mathbb{S}_+^n}(x) + \iota_{[\alpha, \beta]^N}(x)$$

prox_f is intractable.

$$\left. \begin{array}{l} \text{prox}_{f_1} = \text{proj}_{[\alpha, \beta]^N} \\ \text{prox}_{f_2} = \text{proj}_{\mathbb{S}_+^n} \end{array} \right\} \text{known in closed-form.}$$